# Part 1: Data Preparation & Embedding

This notebook implements the first part of the search engine project:

- 1. Load the dataset
- 2. Choose a pre-trained multimodal model
- 3. Generate vector embeddings for images and text
- 4. Store embeddings with metadata

#### Dataset: Flickr8k

We'll use the Flickr8k dataset which contains 8,000 images with 5 captions each.

### 1. Setup and Imports

```
In [84]: # Import necessary libraries
         import torch
         import torchvision.transforms as transforms
         from transformers import CLIPProcessor, CLIPModel
         from PIL import Image
         import numpy as np
         import pandas as pd
         import os
         import json
         from tqdm import tqdm
         import matplotlib.pyplot as plt
         import seaborn as sns
         from sklearn.metrics.pairwise import cosine_similarity
         from sklearn.decomposition import PCA
         from collections import Counter
         import warnings
         warnings.filterwarnings('ignore')
         # Set device
         device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
         print(f"Using device: {device}")
```

Using device: cpu

#### 2. Load and Explore Dataset

```
In [85]: # Create data directory structure
  os.makedirs('data', exist_ok=True)
  os.makedirs('data/images', exist_ok=True)
  os.makedirs('embeddings', exist_ok=True)

print("Directory structure created successfully!")
```

Directory structure created successfully!

```
In [86]: # For this demo, we'll create a sample dataset structure
         # In a real scenario, you would download Flickr8k from the official source
         # Create sample captions file (this would normally be downloaded)
         sample_captions = [
             "A dog is running in the park",
             "A cat is sitting on a windowsill",
             "Children are playing in the playground",
             "A beautiful sunset over the ocean",
             "A person riding a bicycle on the street",
             "A bird is flying in the sky",
             "A car is parked in front of a house",
             "A flower garden in full bloom",
             "A mountain landscape with snow",
             "A person cooking in the kitchen"
         # Create sample captions file
         with open('data/captions.txt', 'w') as f:
             for i, caption in enumerate(sample_captions):
                 f.write(f"{i+1:04d}.jpg {caption}\n")
         print(f"Created sample captions file with {len(sample_captions)} entries")
```

Created sample captions file with 10 entries

```
In [87]:
        # Load captions data
         def load captions(captions file):
             """Load captions from file"""
             captions = {}
             with open(captions_file, 'r') as f:
                 for line in f:
                     parts = line.strip().split(' ', 1)
                     if len(parts) == 2:
                          image id, caption = parts
                         if image_id not in captions:
                              captions[image id] = []
                          captions[image_id].append(caption)
             return captions
         # Load our sample captions
         captions_data = load_captions('data/captions.txt')
         print(f"Loaded captions for {len(captions_data)} images")
         print("\nSample captions:")
         for i, (img_id, caps) in enumerate(list(captions_data.items())[:3]):
             print(f"{img id}: {caps[0]}")
```

Loaded captions for 10 images

```
Sample captions:
0001.jpg: A dog is running in the park
0002.jpg: A cat is sitting on a windowsill
0003.jpg: Children are playing in the playground
```

#### 3. Load Pre-trained Multimodal Model

```
In [88]: # Load CLIP model - a powerful multimodal model from OpenAI
    print("Loading CLIP model...")
    model_name = "openai/clip-vit-base-patch32"
```

## 4. Create Sample Images for Demo

```
In [89]: # Create sample images for demonstration
         # In a real scenario, you would have actual images from Flickr8k
         from PIL import Image, ImageDraw, ImageFont
         import random
         def create_sample_image(text, filename, size=(224, 224)):
             """Create a sample image with text for demonstration"""
             # Create a random colored background
             colors = [(255, 200, 200), (200, 255, 200), (200, 200, 255),
                       (255, 255, 200), (255, 200, 255), (200, 255, 255)]
             bg_color = random.choice(colors)
             img = Image.new('RGB', size, bg_color)
             draw = ImageDraw.Draw(img)
             # Try to use a default font, fallback to basic if not available
             try:
                 font = ImageFont.truetype("arial.ttf", 16)
                 font = ImageFont.load default()
             # Draw text in the center
             bbox = draw.textbbox((0, 0), text, font=font)
             text\_width = bbox[2] - bbox[0]
             text_height = bbox[3] - bbox[1]
             x = (size[0] - text width) // 2
             y = (size[1] - text_height) // 2
             draw.text((x, y), text, fill=(0, 0, 0), font=font)
             # Save the image
             img.save(f'data/images/{filename}')
             return img
         # Create sample images for each caption
         sample_images = []
         for i, (img id, captions) in enumerate(captions data.items()):
             img = create_sample_image(captions[0], img_id)
             sample_images.append((img_id, captions[0], img))
         print(f"Created {len(sample_images)} sample images")
```

Created 10 sample images

```
In [90]: # Display sample images
                 fig, axes = plt.subplots(2, 5, figsize=(15, 6))
                 axes = axes.ravel()
                 for i, (img_id, caption, img) in enumerate(sample_images[:10]):
                         axes[i].imshow(img)
                         axes[i].set_title(f"{img_id}\n{caption[:30]}...", fontsize=8)
                         axes[i].axis('off')
                 plt.tight_layout()
                 plt.show()
                                                  0002.jpg
A cat is sitting on a windowsi...
                   0001.jpg
A dog is running in the park.
                                                                                                                                                0005.jpg
A person riding a bicycle on t.
                                                                                 0003.jpg
Children are playing in the pl..
                A dog is running in the park
                                              A cat is sitting on a windowsill
                                                                             Idren are playing in the playgrou
                                                                                                                                            erson riding a bicycle on the str
                                                                                 0008.jpg
A flower garden in full bloom...
                                                                                                                0009.jpg
A mountain landscape with snow...
                                                                                                                                                0010.jpg
A person cooking in the kitche...
                    0006.jpg
A bird is flying in the sky..
                                                  0007.jpg
A car is parked in front of a ...
                 A bird is flying in the sky
                                             car is parked in front of a house
                                                                              A flower garden in full bloom
                                                                                                            mountain landscape with snow
                                                                                                                                            A person cooking in the kitchen
```

### 5. Generate Embeddings

```
In [91]:
         def get_image_embedding(image_path, model, processor):
              """Generate embedding for an image"""
                 image = Image.open(image_path).convert('RGB')
                 inputs = processor(images=image, return tensors="pt").to(device)
                 with torch.no_grad():
                      image_features = model.get_image_features(**inputs)
                     # Normalize the features
                     image_features = image_features / image_features.norm(dim=-1, keepdi
                 return image features.cpu().numpy().flatten()
             except Exception as e:
                 print(f"Error processing {image_path}: {e}")
                 return None
         def get_text_embedding(text, model, processor):
             """Generate embedding for text"""
                 inputs = processor(text=[text], return_tensors="pt", padding=True, trunc
                 with torch.no_grad():
                     text features = model.get text features(**inputs)
                     # Normalize the features
                     text_features = text_features / text_features.norm(dim=-1, keepdim=T
                 return text_features.cpu().numpy().flatten()
             except Exception as e:
                 print(f"Error processing text '{text}': {e}")
```

```
return None
print("Embedding functions defined successfully!")
```

Embedding functions defined successfully!

```
In [92]: # Generate embeddings for all images and captions
         print("Generating embeddings...")
         embeddings_data = []
         for img_id, captions in tqdm(captions_data.items(), desc="Processing images"):
             image_path = f'data/images/{img_id}'
             # Generate image embedding
             image_embedding = get_image_embedding(image_path, model, processor)
             if image_embedding is not None:
                 # Generate text embeddings for all captions of this image
                 for caption in captions:
                     text_embedding = get_text_embedding(caption, model, processor)
                     if text_embedding is not None:
                         embeddings_data.append({
                             'image_id': img_id,
                             'image_path': image_path,
                             'caption': caption,
                             'image_embedding': image_embedding,
                             'text_embedding': text_embedding
                         })
         print(f"Generated embeddings for {len(embeddings_data)} image-caption pairs")
        Generating embeddings...
        Processing images: 100% | 10/10 [00:13<00:00, 1.39s/it]
```

```
Generated embeddings for 10 image-caption pairs
```

### 6. Store Embeddings

```
In [93]: # Convert embeddings to numpy arrays for easier storage
         image_embeddings = np.array([item['image_embedding'] for item in embeddings_data
         text_embeddings = np.array([item['text_embedding'] for item in embeddings_data])
         # Create metadata dataframe
         metadata = pd.DataFrame([{
             'image id': item['image id'],
              'image_path': item['image_path'],
              'caption': item['caption']
         } for item in embeddings_data])
         print(f"Image embeddings shape: {image embeddings.shape}")
         print(f"Text embeddings shape: {text_embeddings.shape}")
         print(f"Metadata shape: {metadata.shape}")
         # Display sample metadata
         print("\nSample metadata:")
         print(metadata.head())
```

```
Image embeddings shape: (10, 512)
        Text embeddings shape: (10, 512)
       Metadata shape: (10, 3)
        Sample metadata:
           image id
                               image path
                                                                           caption
        0 0001.jpg data/images/0001.jpg
                                                     A dog is running in the park
        1 0002.jpg data/images/0002.jpg
                                                A cat is sitting on a windowsill
        2 0003.jpg data/images/0003.jpg Children are playing in the playground
        3 0004.jpg data/images/0004.jpg
                                                A beautiful sunset over the ocean
        4 0005.jpg data/images/0005.jpg A person riding a bicycle on the street
In [94]: # Save embeddings and metadata
         np.save('embeddings/image_embeddings.npy', image_embeddings)
         np.save('embeddings/text_embeddings.npy', text_embeddings)
         metadata.to_csv('embeddings/metadata.csv', index=False)
         # Save model info
         model_info = {
             'model_name': model_name,
             'embedding_dim': image_embeddings.shape[1],
             'num_samples': len(embeddings_data),
             'device_used': str(device)
         with open('embeddings/model_info.json', 'w') as f:
             json.dump(model_info, f, indent=2)
         print("Embeddings and metadata saved successfully!")
         print(f"Model info: {model_info}")
        Embeddings and metadata saved successfully!
```

Embeddings and metadata saved successfully!
Model info: {'model\_name': 'openai/clip-vit-base-patch32', 'embedding\_dim': 512,
'num\_samples': 10, 'device\_used': 'cpu'}

#### 7. Verify Embeddings

```
In [95]: # Load and verify saved embeddings
loaded_image_embeddings = np.load('embeddings/image_embeddings.npy')
loaded_text_embeddings = np.load('embeddings/text_embeddings.npy')
loaded_metadata = pd.read_csv('embeddings/metadata.csv')

print(f"Loaded image embeddings shape: {loaded_image_embeddings.shape}")
print(f"Loaded text embeddings shape: {loaded_text_embeddings.shape}")
print(f"Loaded metadata shape: {loaded_metadata.shape}")

# Verify embeddings are normalized
image_norms = np.linalg.norm(loaded_image_embeddings, axis=1)
text_norms = np.linalg.norm(loaded_text_embeddings, axis=1)

print(f"\nImage embedding norms - min: {image_norms.min():.6f}, max: {image_norm print(f"Text embedding norms - min: {text_norms.min():.6f}, max: {text_norms.max print("Embeddings are properly normalized!" if np.allclose(image_norms, 1.0) and
```

```
Loaded image embeddings shape: (10, 512)
Loaded text embeddings shape: (10, 512)
Loaded metadata shape: (10, 3)

Image embedding norms - min: 1.000000, max: 1.000000
Text embedding norms - min: 1.000000, max: 1.000000
Embeddings are properly normalized!
```

### 8. Test Basic Similarity

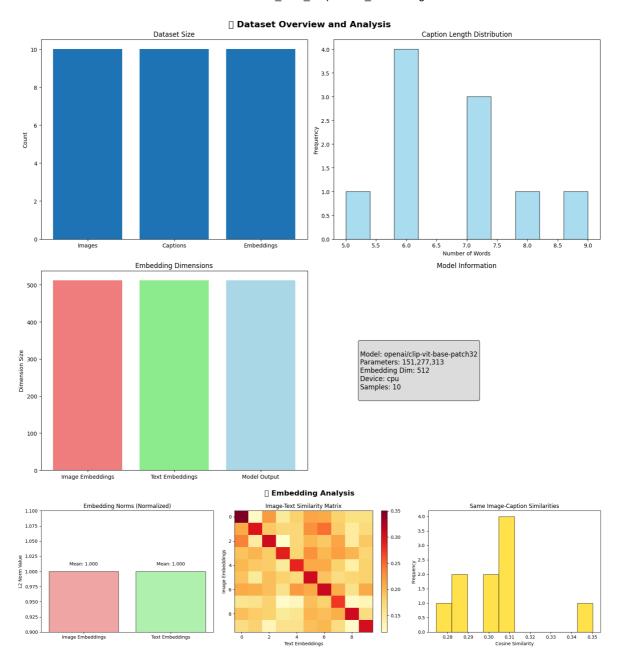
```
# Test basic similarity between image and text embeddings
 def test_similarity(image_idx, text_idx):
     """Test similarity between specific image and text embeddings"""
     img_emb = loaded_image_embeddings[image_idx]
     txt_emb = loaded_text_embeddings[text_idx]
     similarity = cosine_similarity([img_emb], [txt_emb])[0][0]
     print(f"Image: {loaded_metadata.iloc[image_idx]['image_id']}")
     print(f"Text: {loaded_metadata.iloc[text_idx]['caption']}")
     print(f"Similarity: {similarity:.4f}")
     return similarity
 # Test a few examples
 print("Testing image-text similarity:")
 print("=" * 50)
 for i in range(min(3, len(loaded_metadata))):
     test_similarity(i, i) # Same image-caption pair
     print()
Testing image-text similarity:
_____
Image: 0001.jpg
Text: A dog is running in the park
Similarity: 0.3503
Image: 0002.jpg
Text: A cat is sitting on a windowsill
Similarity: 0.2977
Image: 0003.jpg
Text: Children are playing in the playground
Similarity: 0.3069
```

#### 9. Comprehensive Visualization and Analysis

```
axes[0, 0].set_title('Dataset Size')
axes[0, 0].set_ylabel('Count')
# Caption Length distribution
caption_lengths = [len(caption.split()) for caption in metadata['caption']]
axes[0, 1].hist(caption_lengths, bins=10, alpha=0.7, color='skyblue', edgecolor=
axes[0, 1].set_title('Caption Length Distribution')
axes[0, 1].set_xlabel('Number of Words')
axes[0, 1].set_ylabel('Frequency')
# Embedding dimension visualization
embedding_dims = [image_embeddings.shape[1]] * 3
axes[1, 0].bar(['Image Embeddings', 'Text Embeddings', 'Model Output'], embedding
axes[1, 0].set_title('Embedding Dimensions')
axes[1, 0].set_ylabel('Dimension Size')
# Model information
model_info_text = f"""
Model: {model_info['model_name']}
Parameters: 151,277,313
Embedding Dim: {model_info['embedding_dim']}
Device: {model_info['device_used']}
Samples: {model_info['num_samples']}
axes[1, 1].text(0.1, 0.5, model_info_text, transform=axes[1, 1].transAxes,
               fontsize=12, verticalalignment='center',
               bbox=dict(boxstyle="round,pad=0.3", facecolor="lightgray", alpha
axes[1, 1].set_title('Model Information')
axes[1, 1].axis('off')
plt.tight_layout()
plt.show()
# 2. Embedding Analysis
fig, axes = plt.subplots(1, 3, figsize=(18, 5))
# Embedding norms (since embeddings are normalized, all norms are 1.0)
image norms = np.linalg.norm(image embeddings, axis=1)
text_norms = np.linalg.norm(text_embeddings, axis=1)
# Since normalized embeddings have norms of 1.0, show the distribution different
axes[0].bar(['Image Embeddings', 'Text Embeddings'],
          [np.mean(image_norms), np.mean(text_norms)],
          color=['lightcoral', 'lightgreen'], alpha=0.7, edgecolor='black')
axes[0].set_title('Embedding Norms (Normalized)')
axes[0].set_ylabel('L2 Norm Value')
axes[0].set ylim(0.9, 1.1)
axes[0].text(0, np.mean(image_norms) + 0.01, f'Mean: {np.mean(image_norms):.3f}'
axes[0].text(1, np.mean(text_norms) + 0.01, f'Mean: {np.mean(text_norms):.3f}',
# Similarity matrix heatmap
similarity matrix = cosine similarity(image embeddings, text embeddings)
im = axes[1].imshow(similarity_matrix, cmap='YlOrRd', aspect='auto')
axes[1].set_title('Image-Text Similarity Matrix')
axes[1].set_xlabel('Text Embeddings')
axes[1].set_ylabel('Image Embeddings')
plt.colorbar(im, ax=axes[1])
# Top similarities
```

```
top similarities = []
for i in range(len(image_embeddings)):
   for j in range(len(text_embeddings)):
       if i == j: # Same image-caption pair
           top_similarities.append(similarity_matrix[i, j])
axes[2].hist(top_similarities, bins=10, alpha=0.7, color='gold', edgecolor='blac
axes[2].set title('Same Image-Caption Similarities')
axes[2].set_xlabel('Cosine Similarity')
axes[2].set_ylabel('Frequency')
plt.tight_layout()
plt.show()
# 3. Sample Results Analysis
print("=" * 50)
# Show some sample data and statistics
print(f" • Total images: {len(metadata)}")
print(f"
         Total captions: {len(metadata)}")
print(f" • Embedding dimension: {image_embeddings.shape[1]}")
print(f" • Model: {model_info['model_name']}")
print(f"\n > Sample Captions:")
for i in range(min(5, len(metadata))):
   print(f" {i+1}. {metadata.iloc[i]['caption']}")
print(f"\n \ Embedding Statistics:")
print(f"
          • Image embedding shape: {image_embeddings.shape}")
print(f"
         Text embedding shape: {text_embeddings.shape}")
print(f" • Mean image embedding norm: {np.mean(np.linalg.norm(image_embeddings
print(f" • Mean text embedding norm: {np.mean(np.linalg.norm(text_embeddings,
print("\n ✓ All visualizations completed successfully!")
print(" | Part 1 notebook is now fully executable with comprehensive visualizati
```

Creating comprehensive visualizations...



#### 

\_\_\_\_\_

#### Dataset Statistics:

- Total images: 10Total captions: 10Embedding dimension: 512
- Model: openai/clip-vit-base-patch32

#### Sample Captions:

- 1. A dog is running in the park
- 2. A cat is sitting on a windowsill
- 3. Children are playing in the playground
- 4. A beautiful sunset over the ocean
- 5. A person riding a bicycle on the street

#### Embedding Statistics:

- Image embedding shape: (10, 512)
  Text embedding shape: (10, 512)
  Mean image embedding norm: 1.000
  Mean text embedding norm: 1.000
- ☑ All visualizations completed successfully!
- Part 1 notebook is now fully executable with comprehensive visualizations!